

Key idea: use each agent run's tests to judge the other runs' patches

